



Grounded Decoding: Guiding Text Generation with Grounded Models for Robot Control

Wenlong Huang^{1*}, Fei Xia², Dhruv Shah³, Danny Driess², Andy Zeng², Yao Lu², Pete Florence², Igor Mordatch⁴, Sergey Levine²³, Karol Hausman² and Brian Ichter²

^{*}Work done as an intern at Google, ¹Stanford University, ²Robotics at Google, ³UC Berkeley, ⁴Google Research

Recent progress in large language models (LLMs) has demonstrated the ability to learn and leverage Internet-scale knowledge through pre-training with autoregressive models. Unfortunately, applying such models to settings with embodied agents, such as robots, is challenging due to their lack of experience with the physical world, inability to parse non-language observations, and ignorance of rewards or safety constraints that robots may require. On the other hand, language-conditioned robotic policies that learn from interaction data can provide the necessary grounding that allows the agent to be correctly situated in the real world, but such policies are limited by the lack of high-level semantic understanding due to the limited breadth of the interaction data available for training them. Thus, if we want to make use of the semantic knowledge in a language model while still situating it in an embodied setting, we must construct an action sequence that is *both* likely according to the language model and also realizable according to grounded models of the environment. We frame this as a problem similar to probabilistic filtering: decode a sequence that both has high probability under the language model and high probability under a set of grounded model objectives. We demonstrate this guided decoding strategy is able to solve complex, long-horizon embodiment tasks in a robotic setting by leveraging the knowledge of both models. The project's website can be found at grounded-decoding.github.io.

1. Introduction

Recent works have demonstrated robots that are increasingly proficient at understanding and acting upon natural language, whether through planning or conditioned policies. Complementing such progress, the field of natural language processing has recently seen large language models (LLMs) become ubiquitously used as pre-trained or few-shot prompted models, due to their impressive few-shot performance and vast knowledge-base. These LLMs have efficiently learned from web-scale data through autoregressively modeling the probability distribu-

tion over text tokens and thus generate text. However, the nature of this process is such that applying such models to embodied settings remains a challenge. They have not interacted with their environment, lack observability of non-language observation modalities (e.g., images), and may not know what is safe or possible for a particular embodiment.

Determining how to execute long-horizon behaviors based on high-level verbal commands is one particular area of robotics where the rich semantic knowledge in large language models can be especially useful. This problem combines elements of

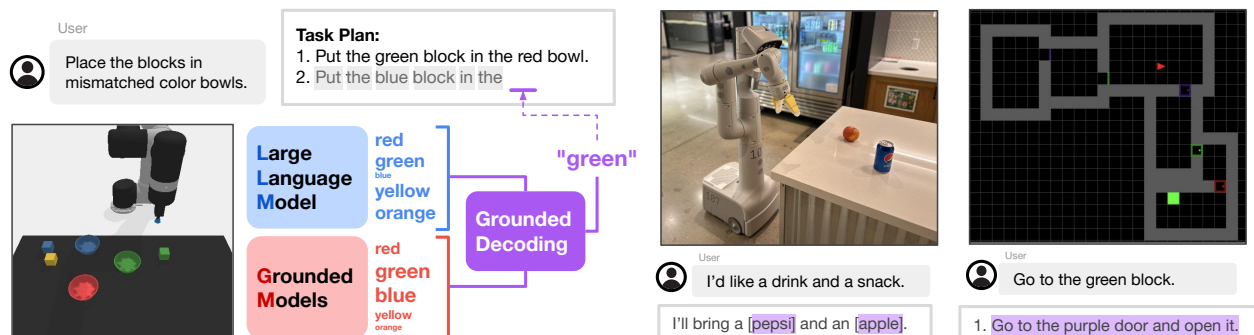


Figure 1: Grounded Decoding solves robotic tasks by taking an instruction as input and selecting tokens that have high probability under a Large Language Model (LLM) and a set of Grounded Models (GM). Thus, it leverages the open-vocabulary and semantic knowledge of LLMs while being grounded in the environment and in the robot's capabilities. Furthermore, the whole process does not require expensive fine-tuning of the LLM.

semantic reasoning and planning: the robot must understand the instruction, determine the steps needed to fulfill it, and also determine how to sequence those steps appropriately given its capabilities and the current state of the environment. However, this is not a problem that can be solved purely with semantics, as it requires sufficient grounding to understand how the task should be performed *in context* – for example, in the example in Figure 1, the language model alone has no way of knowing which block to pick up because this requires knowledge of which blocks are present, and also what manipulations the robot is capable of performing on them. Thus, although a language model can assign probabilities for how likely various steps are to correspond to the desired task *semantically*, the constraints of the planning problem must also enter into the process. These constraints could themselves be represented as probabilities that mirror the token probabilities generated by a language model, reflecting their applicability to the current environment rather than their semantic likelihood. We can frame this as a problem similar to probabilistic filtering: decode a sequence (i.e., a task description) that both has a high probability under the language model and a high probability under a grounded model that predicts how applicable this sequence is to the current scene.

Herein, we present Grounded Decoding (GD), a scalable, general approach to planning with LLMs embodied domains. Grounded Decoding jointly decodes the token probability of an LLM and token probabilities from token-conditioned, robotic functions, such as affordance functions capturing the abilities of a robot given its embodiment, safety functions, or more. By guiding the LLM directly at its output, Grounded Decoding enables a general and flexible family of planning algorithms. We present empirical results on three robotic domains, demonstrating GD’s the necessity of grounding models and the ability of GD to perform long-horizon, complex tasks.

2. Related Work

Guided Decoding for Language Models. Decoding strategies for large language models is an active area of research within natural language processing [1, 2, 3, 4, 5]. A number of recent works have focused on developing decoding heuristics for natural text generation [6, 7, 8, 9, 3, 10, 11, 12]. Another line of works use external classifiers for maximiz-

ing certain language-space utilities when decoding language models [13, 14, 15, 16, 17, 18, 5, 19, 20, 21]. Most closely related to our work are classifier-guided decoding methods developed for offline domains, such as image captioning [22, 23] and task-oriented dialog [24, 25]. However, extensions to embodied domains, which we investigate exclusively in this work, remain non-trivial because grounding in embodied domains is bounded by the abilities of the agent and by environment state transition as the agent actively interacts with the environment.

Embodied and Multimodal Language Models. Training language models to understand embodiment is an active area of research. Training multimodal models can enable some degree of embodied reasoning, such as understanding images and videos [26, 27, 28, 29]. Directly finetuning language models to output actions has also been investigated [30, 31, 32]. Lastly, training downstream models on language model embeddings shows promise [33, 34, 35, 36, 37, 38]. In this work, we investigate leveraging large frozen language models for embodied applications [39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50], with grounded models to provide domain-specific grounding during decoding process.

Comparison to SayCan. The most closely related work to our work is SayCan [40]. SayCan uses a large language model and a value function to select robotic skills among a constrained set of primitives. This constrained set of primitives enables SayCan to use the so-called “scoring-mode” of the LLM to get the probability of a skill being useful to a high-level instruction. This requirement to consider only a fixed and enumerated set of primitives limits the applicability of SayCan in scenarios with many possible skills, such as open vocabulary or combinatorial tasks. Grounded Decoding on the other hand jointly decodes the LLM and the grounded model at the token level, allowing for flexible and expressive decoding of open vocabulary tasks. Furthermore, SayCan considers only grounding functions derived from RL-trained value functions for affordance grounding functions, while Grounded Decoding explores many types of grounding functions to propose a broad family of algorithms.

Task and Motion Planning. Task and motion planning [51] seeks to solve high-level instructions via sequencing tasks in dynamically feasible manner. Research within this area generally focuses on symbolic

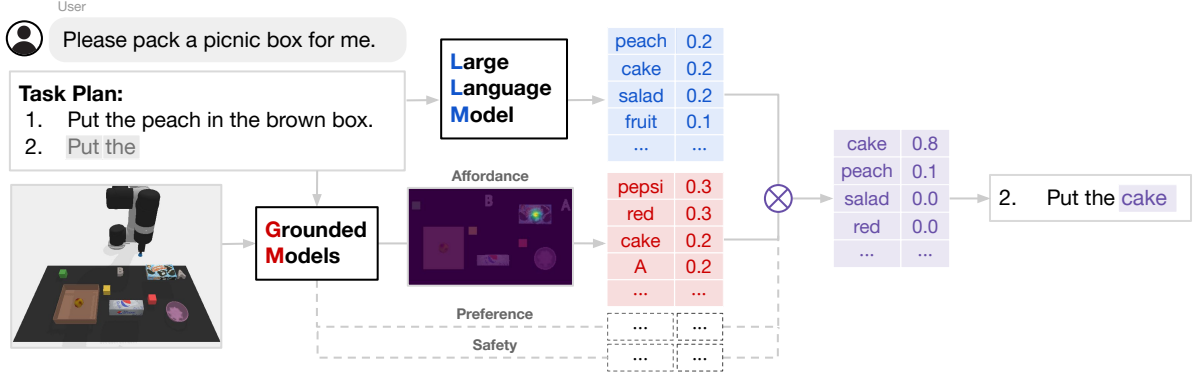


Figure 2: Overview of **Grounded Decoding**. Given a *free-form* language instruction, a **language model** and **grounded models** jointly decide the next candidate token to be decoded by **combining** their respective likelihood. Language model proposes likely tokens that produce goal-directed and coherent long-horizon behaviors, while grounded models connect them to the physical scene, through a flexible composition of multiple objective functions, such as affordance, preferences, and safety, that can be independently obtained.

planning [52] or optimization-based [53] approaches. Machine learning has increasingly been used to accelerate planning and enable new domains [54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67]. However, planning constraints are often explicitly specified for TAMP methods. In contrast, we specify constraints as (learned) probabilities, which are baked into the decoding process and provided by domain-specific grounded models.

3. Grounded Decoding

3.1. LLMs and Grounding Models

Large Language Models. LLMs are trained to predict the probability $p(W)$ of a text sequence W , represented as a sequence of tokens $W = w_{1:N} = (w_1, \dots, w_N)$. The tokens are elements of a fixed vocabulary \mathcal{W} . Typical neural architectures factorize the joint probability into $p(W) = \prod_{n=1}^N p_{\text{LLM}}(w_n | w_{1:n-1})$, where p_{LLM} is predicted by a transformer network [68]. Given p_{LLM} , generating a text consisting of N -many tokens, the so-called decoding process, can be seen as the optimization problem $\arg \max_{w_{1:N} \in \mathcal{W}} \prod_{n=1}^N p_{\text{LLM}}(w_n | w_{1:n-1})$, which in practice is solved, e.g., using greedy search ($w_n = \arg \max_{w \in \mathcal{W}} p_{\text{LLM}}(w | w_{1:n-1})$), beam search, or sampling strategies. To further ensure the LLM is solving a desired task, one typically starts with a given text, the so-called *prefix* or *prompt*, that describes the task, and then the LLM completes this task in its decoding process.

Grounding Functions. We use the concept of grounding functions, $p_G(w_{1:n} | s)$, which seek to model a probability of tokens $w_{1:n}$ given state $s \in \mathcal{S}$. This state is intended to capture the embodiment

of the robot and the environment, and may be for example an image, a proprioceptive robot state, or the environment state. Thus the grounding function models probabilities relevant to the robot embodiment and environment, such as whether the tokens are possible for the robot to execute given the state (affordances), or other values like safety, cost, or user preferences.

3.2. Problem formulation.

Given an instruction in language ℓ , we can use an LLM to produce a language plan $w_{1:N}$ by finding the most likely decoding according to the LLM’s probability distribution. $p_{\text{LLM}}(w_{1:N} | \ell)$, with ℓ being the prefix. However, based on the instruction ℓ as the prefix alone, the LLM can easily generate text that is not grounded in the physical state of the environment, rendering such plans useless in the real world. In order to *ground* the language model in an actual physical embodiment, we propose *Grounded Decoding* (GD): The main idea of GD is to guide the generation of token sequences with a *grounding function* that is conditioned on the embodiment of the system.

Formally, let $s \in \mathcal{S}$ denote a representation of the state of the world (which could be given by an image). Then, GD

$$w_{1:N}^* = \arg \max_{w_{1:N}, w_n \in \mathcal{W}} p_{\text{GD}}(w_{1:N} | s, \ell) \quad (1)$$

attempts to generate text that is consistent with both the instruction ℓ and the physical state s .

To leverage the knowledge present in LLMs from internet-scale data, we factorize $p_{\text{GD}}(w_{1:N} | s, \ell)$ as

follows

$$p_{\text{GD}}(w_{1:N}|s, \ell) = \frac{p(s, \ell|w_{1:N}) p(w_{1:N})}{p(s, \ell)} \quad (2)$$

$$= \frac{p(s|w_{1:N})p(\ell|w_{1:N})p(w_{1:N})}{p(s, \ell)} \quad (3)$$

$$= \frac{p(w_{1:N}|\ell)p(\ell)p(w_{1:N}|s)p(s)p(w_{1:N})}{p(s, \ell)p(w_{1:N})p(w_{1:N})} \quad (4)$$

$$\propto \frac{p(w_{1:N}|\ell)}{p(w_{1:N})} p(w_{1:N}|s) \quad (5)$$

$$\propto p(w_{1:N}|\ell)p(w_{1:N}|s). \quad (6)$$

We make two assumptions for this derivation, that s and ℓ are conditionally independent given $w_{1:N}$ (required for Line 5) and that $p(w_{1:N})$ is uniform over responses (required for Line 6). We then factorize into token decoding

$$p_{\text{GD}}(w_{1:N}|s, \ell) \propto \prod_{n=1}^N p_{\text{LLM}}(w_n|w_{1:n-1}, \ell) p_{\text{G}}(w_{1:n}|s). \quad (7)$$

The first term, $p_{\text{LLM}}(w_n|w_{1:n-1}, \ell)$, can be modeled as the probability of the LLM predicting the token for the given instruction ℓ appended previously decoded tokens $w_{1:n-1}$ without the state s as input. The second term, $p_{\text{G}}(w_{1:n}|s)$, is the grounding function that is only conditioned on the state s and judges whether the generated text $w_{1:n}$ is consistent with the physical state. The core idea behind this factorization is that LLMs exhibit long-term planning capabilities, while the grounding function guides the planning of the LLM to be possible in the concrete embodied physical world without needing to be informed or capable of reasoning over the long-horizon instruction.

3.3. Algorithm – Grounded Decoding.

Algorithm 1 details the steps in GD and Figure 2 visualizes a single step of the simplest greedy search form of GD. GD proceeds through a process similar to probabilistic filtering by selecting tokens iteratively that have high probability under the language model and the grounded model. After each token is selected, it is appended to the prefix. The process continues until a token in the terminal set $\mathcal{W}_{\text{term}}$ is selected, at which point the command is returned and executed. Additionally, we note that GD, in its essence, provides a grounded scoring function; thus, it can be easily extended to any search methods such as beam search, top-k sampling, etc.

Algorithm 1 GD w/ Greedy Search

- 1: **Given:** state s , instruction ℓ , terminal set $\mathcal{W}_{\text{term}}$
 - 2: **Initialize:** $w = \{\}$, $n = 0$
 - 3: **while** $w_n \notin \mathcal{W}_{\text{term}}$ **do**
 - 4: $n = n + 1$
 - 5: $w_n = \arg \max_{w_n \in \mathcal{W}} p_{\text{LLM}}(w_n|w_{1:n-1}, \ell) p_{\text{G}}(w_{1:n}|s)$
 - 6: **Return:** w
-

Grounded Decoding for Robot Control. This work investigates grounded decoding in the context of robot control. For this setting, we assume to have access to a language conditioned policy π that predicts actions $a = \pi(s, w_{1:i})$ based on a language command $w_{1:i}$ and state s . We further assume access to a grounding function $p_{\text{G}}(w_{1:n}|s)$ that model embodied probabilities, e.g., affordances or safety of a the command given the state. Crucially, this grounding function must accept partial commands to enable grounding during decoding. As an example, an affordance ground function for a skill “pick up *object*”, should emit a high probability for “pick” and “pick up” if any object is able to be picked and collapse to only feasible objects only once the *object* token is decoded. Once a complete, executable skill is fully decoded it is executed in the environment and the decoding continues with the new state updated for the grounding function.

4. Experiments

Our experiments section demonstrates Grounded Decoding on three different environments (Figure 3) and with a variety of grounding functions in order to demonstrate its generality and flexibility. We show first a tabletop manipulation environment with grounding functions for affordances, safety, and preferences. Second, we show a 2D Maze environment built from Minigrid [69], with RL-trained value function-based grounding functions. Finally, we show a real robot in an office kitchen, inspired by [40], with a CLIP-based grounding function. We provide high-level analysis of the results in this section and detailed analysis in Section 5.

4.1. Long-Horizon Tabletop Manipulation

Herein we experiment with a simulated tabletop manipulation scene based on the RAVENS [70] environment and a CLIPort [71] policy. We create a custom set of 20 tasks, all of which are specified via natural language instructions and require the agent to



Figure 3: Example rollouts and likelihood of representative tokens under Grounded Decoding objective in three distinct domains: simulated tabletop rearrangement (*top*), Minigrid 2D Maze (*middle*), and real-world kitchen mobile manipulation (*bottom*). Each domain uses different prompts, grounded models, and low-level primitives. The GD formulation is shared across the domains, decoding a pre-trained language model with respect to domain-specific grounded models to decompose a *open-ended* instruction into actionable steps.

perform long-horizon reasoning based on its observation of the environment. The tasks are split into 10 seen tasks and 10 unseen tasks, where seen tasks may be used for training (in the case of supervised baseline) or used as few-shot prompting. They are grouped by three categories:

- i. **Letters (8 tasks):** Rearranging alphabetical letters (e.g., “put the letters in alphabetical order from left to right”).
- ii. **Blocks & Bowls (8 tasks):** Rearranging or combining blocks and bowls (e.g., “put the blocks in the matching bowls”).
- iii. **Box Packing (4 tasks):** Sorting food items and utensils into boxes in accordance with safety constraints and user preferences (e.g., “Can you pack the picnic box for me?”).

Given only high-level language instructions and top-down visual observation of the environment, Grounded Decoding uses language models as planners to decompose the high-level instructions into a sequence of steps [39, 40]. After a complete step command is generated, we use a pre-trained multi-task language-conditioned CLIPort [71] as the primitive policy to execute the command. Note that because GD generated grounded *free-form* actions, it does not require each step to strictly map to a repertoire of skill as in [39, 40].

While the GD formulation is amenable for any decoding method, we study two variants using beam search and greedy search. For baselines, we compare to “No Grounding” baseline that decodes only according to language model likelihood. Furthermore, we compare to solitary method CLIPort [71] that directly take in the high-level language instructions without a planner. We consider two variants of CLIPort: 1) “Short” that is trained with only single-step pick-and-place commands, and 2) “Long” that is trained on high-level instructions from the 10 training tasks. For more details, please refer to Appendix A.2.4.

Owing to the elegance and extensibility of Grounded Decoding, we can seamlessly compose multiple grounding functions depending on the task at inference time. To compose grounding functions $\{p_1, p_2, \dots, p_n\}$, the overall grounding score is calculated as $p_G = \prod_{i=1}^n p_i$. Please refer to the Appendix A.1 for which grounding function are used for each task and Appendix A.2.4.

Affordance Grounding Function (AF). The affordance should be a function of the robot, the environment, and the underlying policy in order to accurately determine what is possible in the scene. As part of its architecture, given a scene image s and partially-decoded open-ended language instruction




Figure 4: Example decision-making with GD, where key decoded tokens are shown (*yellow, purple, red, yellow*). For each of the key token position, we show the top 10 candidates with their likelihoods under the **language model**, **grounded model**, and **both**. Combined scores are normalized to the maximum for visual clarity; others are normalized to the sum for showing their relative confidence. While only greedy search is illustrated for simplicity, alternative search methods may achieve better results, as shown in Section 4.

ℓ the CLIPort primitive policy predicts unnormalized logits over the pixel space, which we use directly here as affordances. This grounding is used for all tasks.

Safety Grounding Function (S). World grounding for embodiment need not only be affordances, it can also be crucial robotics functions like safety. While any safety function that can score a language instruction ℓ may be used, we implement a simple indicator to prevent the robot from interacting hazardous objects. This safety grounding function is used for 3 tasks in Box Packing task family.

Preference Grounding Function (P). Robots operating alongside humans should also be aware of human preferences, which often differ based on the specific user. We choose two random objects as the preferred objects to be used for the household object sorting tasks. Note that unlike safety functions, preferences often come in the form of “soft requirement”, and thus we implement them as a bias towards these objects of interest. More generic preference functions may also be learned or statistically calculated based on history of user data. A preference grounding function is used for 2 tasks in Box Packing task family.

Analysis. Results grouped by each task category are shown in Table 1. Please refer to the Appendix for detailed breakdown. Each method is evaluated on 20 episodes for each task within each task category. Supervised methods, such as CLIPort, are found to perform poorly on unseen tasks. Methods that leverage language model planner show better

	CLIPort		+LLM		+GD
	Short	Long	Ungrounded	Greedy	Beam
Seen Tasks					
Letters	7%	40%	20%	43%	57%
Blocks & Bowls	2%	62%	35%	60%	77%
Box Packing*	15%	28%	11%	79 %	78%
Unseen Tasks					
Letters	6%	10%	19%	37%	41%
Blocks & Bowls	6%	10%	28%	44%	50%

Table 1: Average success rate for each tabletop task category. *Box Packing tasks are seen during training, but safety and preference requirements are only enforced during evaluation.

generalization to unseen tasks but fall short due to lack of grounding. Grounded Decoding achieves the best results by enabling the LLM to plan actions using grounded information and is further improved with beam search.

4.2. 2D Maze

We further evaluate the long-horizon reasoning performance of Grounded Decoding for 2D maze-solving on Minigrid tasks [69]. The agent receives a top-down view of the environment along with a natural language instruction. The tasks are grouped in three categories:

- i. **Easy:** Simple tasks where the horizon is short (10-30 steps) and fully described by the textual instruction, e.g. OpenDoors and PutNext.
- ii. **Medium:** Short and long-horizon tasks (up to

80 steps) with step-by-step textual instructions, e.g. `LockedRoom`.

iii. **Hard:** Complex, long-horizon instructions (over 100 steps) with ambiguous instructions that necessitate multi-step reasoning and efficient exploration, e.g. `MultiRoom` and `BlockedUnlock`.

Given such instructions and top-down observations, Grounded Decoding uses an LLM as a planner to decompose the instructions into a sequence of steps. Following the recipe from Section 3.3, we combine this language model planning with an affordance function grounded in the agent’s observations. To obtain this affordance function, we use the value function from the goal-conditioned policy that is trained with PPO [72]. This policy performs short-horizon skills such as “Go to red key” or “Open the door” and are conditioned on CLIP embeddings of the skill and an image of the scene.

Similar to the tabletop manipulation experiments, we compare the two variants of Grounded Decoding— with greedy and beam search — with the solitary policy [72], a hierarchical algorithm that uses the low-level skills and plans over it using model-free RL, and a hierarchical algorithm that uses an ungrounded language model for planning for textual skills [39]. Each method is evaluated with 100 randomly initialized environment layouts and object positions.

	+Skills		+LLM	+GD	
	PPO	HRL	Ungrounded	Greedy	Beam
Easy	28%	68%	96%	100%	100%
Medium	13%	48%	87%	93%	97%
Hard	6%	31%	54%	78%	88%

Table 2: 2D Maze success rates.

Analysis. Table 2 reports the success rate. The “flat” RL agent performs poorly in all but the simplest environments, owing to difficulties with understanding the high-level instruction and reasoning over long horizons (often over 100 steps). Planning over low-level skills using hierarchical RL [73] improves this performance, since the high-level decision-making problem is greatly simplified. However, the high-level RL agent still needs to reason over low-level (textual) skills by understanding their underlying capabilities and stitching them together. Using the planning capabilities of large language models to reason over textual skills significantly

boosts this performance [39], since the language model can inherit the strong reasoning capabilities from its training data. This tends to be insufficient in challenging environments, however, since the number of *potentially viable* skills may be very large and the LLM has no information about the robot’s observations. GD can leverage the learned affordance function (in this case, the goal-conditioned value function) to inform the language model’s plans, enabling successful long-horizon reasoning. We further find that beam search improves performance modestly, particularly in long-horizon tasks.

4.3. Mobile Manipulation in a Kitchen

Our last environment is a kitchen robot in the real world, and we follow the same implementations of the mobile manipulation platform and skills in SayCan [40]. We perform instruction following tasks, as in [40]. An example task is “Bring an apple”, for which the robot needs to plan and execute a sequence of “1. Find an apple, 2. Pick up the apple, 3. Bring it to you. 4. Put it down, 5. Done”. We split the tasks into two categories. *Unambiguous* means the instruction explicitly contains the object of interest, and *Ambiguous* means the instruction does not contain the object name. For example, when human asks “bring me the fruit”, the robot needs to first figure out what fruit is available. We assume when the robot starts, it is able to see all the necessary objects.

Partial GD and Chain-of-thought We make a slight modification of the SayCan [40] algorithm to enable grounded decoding. The language model in SayCan takes in a prompt and a human instruction, and then predicts next steps. When chain of thought prompting [74] is enabled, SayCan algorithm will first decode a chain of thought in generative mode, then proceed with scoring mode to generate the plan. However, this open vocabulary generative process is not amenable to SayCan’s fixed grammar, which can result in errors. We minimally augment the SayCan algorithm to prompt the language model to generate chain-of-thought that is visually grounded. We make sure that objects mentioned in chain of thought are in square bracket, which indicate to the language model that the object herein is “visually grounded”. During the chain-of-thought decoding, the left bracket will trigger grounded decoding and the right bracket will revert to ungrounded decoding. We take in current image as additional input, and use object detection score as grounding function [42]. This is illustrated

in Fig. 5.

Tasks	GD (ours)		SayCan	
	Planning	Execution	Planning	Execution
Unambiguous	85%	57%	85%	57%
Ambiguous	58%	44%	33%	25%

Table 3: Kitchen task planning and execution success rates.

Analysis. We found in Table 3 that we can recover the performance when the queries are explicit, and gain 25% in planning performance when the queries are ambiguous.

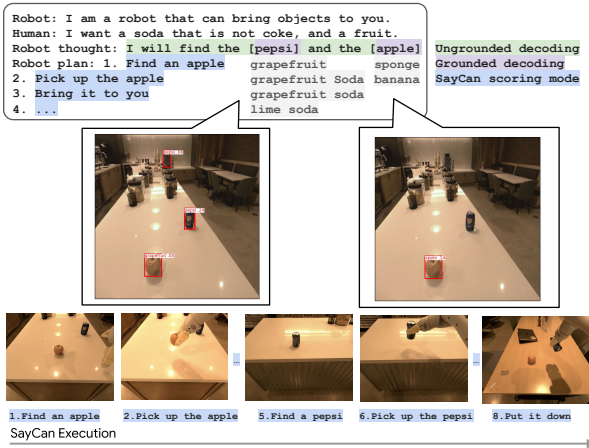


Figure 5: Grounded Decoding with open-vocabulary detection model in the loop resolves ambiguous queries.

5. Analysis

5.1. Comparison to SayCan

In this section, we directly compare GD to SayCan [40], which is related to our method in that both combine language model knowledge and grounded model knowledge (and discussed in more detail in Related Work Section 2). However, SayCan uses the language model to score all pre-specified options, rendering it inefficient at dealing with large or combinatorial action spaces. In contrast, GD computation considers all possible language token in the autoregressive decoding process, which is independent of the size of the action space. Results shown in Table 4 demonstrate that GD is two orders of magnitude more efficient on our tasks, with comparable performance. Furthermore, by decoding at the most basic functioning unit of language, GD’s formulation allows open-vocabulary grounding beyond just

affordances, e.g. safety, preferences, and multimodal embeddings such as CLIP.

	GD (Greedy)	GD (Beam)	SayCan
Success Rate	50%	60%	64%
Token Count	1x	4.3x	113.7x

Table 4: By avoiding full enumeration of the skills, GD is much more efficient than SayCan while staying performant.

5.2. Breakdown of Failure Reasons

Because all hierarchical approaches share an imperfect low-level policy for step execution, the results reported in Table 1 are compounded with both planning failures and low-level policy failure. In Figure 6, we provide failure breakdown analysis for Grounded Decoding and associated baselines. Note that the CLIPort baselines are solitary methods that do not use a planner, so the failures are solely composed of policy failures. As shown in Figure 6, while all planning-based methods use the same underlying low-level policy, Grounded Decoding significantly reduces planning failure by being able to incorporate grounded scene information into the decoding process. Moreover, we observe that despite the shared affordance function across beam and greedy search, the beam search variant performs stronger by being aware of the full-length single-step instructions during decoding. A notable planning failure occurs when the planner produces a slightly incorrect plan (“orange block” instead of “yellow block”), but as the grounding and the policy share limitations, the policy often still interacts with the correct object.

5.3. Grounded Action Manifold

A central goal of this work is to investigate the integration of grounded information into language model decoding to output instructions actionable by a policy. To investigate this, we use a t-SNE [75] plot to illustrate the extent to which grounded models help narrow down the search space for language models. Specifically, we first enumerate all meaningful instructions in the tabletop domain, such as “pick x and place it on y,” which are represented as dots in the figure. We then compute the affordance values with respect to four different scenes, where each color represents one scene. Finally, we group the dots using t-SNE and BERT embeddings. The figure shows that grounded models can effectively




Figure 6: Failure breakdown analysis of various baselines showing that while all planning-based methods use the same underlying low-level policy, GD significantly reduces planning failure by incorporating grounded scene information into the decoding process. The beam search variant outperforms greedy search by being aware of the full instruction during decoding.

identify achievable skills to produce an actionable manifold within the language space and that this grounding is required, as language alone does not perfectly group actionable skills. It is worth noting that while we provide manual enumeration of all possible skills for practical analysis, the language space is much larger, composed of roughly a 40,000 token vocabulary at each step. This highlights the even more pronounced narrowing of the search in the language space.




Figure 7: Visualization of actions colored by affordance values in different scenes. Every dot represents a possible action in the tabletop domain, where the majority of the actions are infeasible. We show how grounded models can identify the feasible actions for specific scenes. Notably, these actions are not always clustered in language space, requiring the grounding function to determine what action to perform.

6. Conclusions, Limitations, & Future

We presented Grounded Decoding (GD), an approach for leveraging the knowledge and capabilities of large language models in embodied settings through grounding functions, which model the probabilities

of tokens given an embodiment. GD resembles probabilistic filtering, by decoding tokens that have high probabilities under the language model *and* under the grounding model. By guiding the LLM’s decoding directly at its output, GD is a general, flexible, and expressive approach to embodied tasks. This is demonstrated on three embodied domains, showing GD is capable of solving complex, long-horizon tasks.

Though quite general and flexible, GD has a few limitations. First, though the formulation of grounding functions is general, the quality and availability of general grounding functions is still a bottleneck, particularly compared to the capabilities of LLMs. We hope that recent progress in large-scale robotics models (e.g. [76] and [77]) can remove this bottleneck, and note that the flexibility of GD allows such progress to be straightforwardly leveraged. Second, prompt engineering is often required to steer LLMs to the desired action space (e.g., likely action verbs, likely present objects). Finally, the joint decoding may be limiting compared to a single model capable of both grounding and language reasoning [29, 26].

This work presented a family of algorithms for grounding LLMs in embodiment, for which there are many avenues for future work. The flexibility of the approach enables many other grounding functions and ways to integrate grounding. Furthermore, the development and integration of a foundation model for grounding would improve performance significantly. Finally, though GD’s probabilistic filtering-based approach is quite general, fusing grounding information to the language model *after* each token decoding may be limiting and we plan to investigate how such grounding can be elegantly integrated *during* decoding.

Acknowledgments

The authors would like to acknowledge Pierre Sermanet, Carolina Parada, Jie Tan, Yevgen Chebotar, Vincent Vanhoucke, and Dorsa Sadigh for their feedback and contributions. This work is supported in part by OpenAI academic access program, granted to Wenlong Huang.

References

- [1] I. Sutskever, O. Vinyals, and Q. V. Le, “Sequence to sequence learning with neural networks,” *Advances in neural information processing systems*, vol. 27, 2014.

- [2] Y. Wu, M. Schuster, Z. Chen, Q. V. Le, M. Norouzi, W. Macherey, M. Krikun, Y. Cao, Q. Gao, K. Macherey, *et al.*, “Google’s neural machine translation system: Bridging the gap between human and machine translation,” *arXiv preprint arXiv:1609.08144*, 2016.
- [3] A. Holtzman, J. Buys, L. Du, M. Forbes, and Y. Choi, “The curious case of neural text degeneration,” *arXiv preprint arXiv:1904.09751*, 2019.
- [4] S. Welleck, I. Kulikov, J. Kim, R. Y. Pang, and K. Cho, “Consistency of a recurrent language model with respect to incomplete decoding,” *arXiv preprint arXiv:2002.02492*, 2020.
- [5] R. Leblond, J.-B. Alayrac, L. Sifre, M. Pislari, J.-B. Lespiau, I. Antonoglou, K. Simonyan, and O. Vinyals, “Machine translation decoding beyond beam search,” *arXiv preprint arXiv:2104.05336*, 2021.
- [6] C. Meister, T. Vieira, and R. Cotterell, “If beam search is the answer, what was the question?,” *arXiv preprint arXiv:2010.02650*, 2020.
- [7] J. Kasai, K. Sakaguchi, R. L. Bras, D. Radev, Y. Choi, and N. A. Smith, “Beam decoding with controlled patience,” *arXiv preprint arXiv:2204.05424*, 2022.
- [8] C. Meister, T. Pimentel, G. Wiher, and R. Cotterell, “Typical decoding for natural language generation,” *arXiv preprint arXiv:2202.00666*, 2022.
- [9] A. Fan, M. Lewis, and Y. Dauphin, “Hierarchical neural story generation,” *arXiv preprint arXiv:1805.04833*, 2018.
- [10] S. Basu, G. S. Ramachandran, N. S. Keskar, and L. R. Varshney, “Mirostat: A neural text decoding algorithm that directly controls perplexity,” *arXiv preprint arXiv:2007.14966*, 2020.
- [11] N. S. Keskar, B. McCann, L. R. Varshney, C. Xiong, and R. Socher, “Ctrl: A conditional transformer language model for controllable generation,” *arXiv preprint arXiv:1909.05858*, 2019.
- [12] T. Scialom, P.-A. Dray, S. Lamprier, B. Piwowarski, and J. Staiano, “Discriminative adversarial search for abstractive summarization,” in *International Conference on Machine Learning*, pp. 8555–8564, PMLR, 2020.
- [13] C. Snell, I. Kostrikov, Y. Su, M. Yang, and S. Levine, “Offline rl for natural language generation with implicit language q learning,” *arXiv preprint arXiv:2206.11871*, 2022.
- [14] K. Yang and D. Klein, “Fudge: Controlled text generation with future discriminators,” *arXiv preprint arXiv:2104.05218*, 2021.
- [15] A. Holtzman, J. Buys, M. Forbes, A. Bosselut, D. Golub, and Y. Choi, “Learning to write with cooperative discriminators,” *arXiv preprint arXiv:1805.06087*, 2018.
- [16] J. Li, W. Monroe, and D. Jurafsky, “Learning to decode for future success,” *arXiv preprint arXiv:1701.06549*, 2017.
- [17] B. Krause, A. D. Gotmare, B. McCann, N. S. Keskar, S. Joty, R. Socher, and N. F. Rajani, “Gedi: Generative discriminator guided sequence generation,” *arXiv preprint arXiv:2009.06367*, 2020.
- [18] M. Ghazvininejad, X. Shi, J. Priyadarshi, and K. Knight, “Hafez: an interactive poetry generation system,” in *Proceedings of ACL 2017, System Demonstrations*, pp. 43–48, 2017.
- [19] A. Baheti, A. Ritter, J. Li, and B. Dolan, “Generating more interesting responses in neural conversation models with distributional constraints,” *arXiv preprint arXiv:1809.01215*, 2018.
- [20] S. Dathathri, A. Madotto, J. Lan, J. Hung, E. Frank, P. Molino, J. Yosinski, and R. Liu, “Plug and play language models: A simple approach to controlled text generation,” *arXiv preprint arXiv:1912.02164*, 2019.
- [21] T. Hartvigsen, S. Gabriel, H. Palangi, M. Sap, D. Ray, and E. Kamar, “Toxigen: A large-scale machine-generated dataset for adversarial and implicit hate speech detection,” *arXiv preprint arXiv:2203.09509*, 2022.
- [22] Y. Tewel, Y. Shalev, I. Schwartz, and L. Wolf, “Zero-shot image-to-text generation for visual-semantic arithmetic,” *arXiv preprint arXiv:2111.14447*, 2021.

- [23] Y. Su, T. Lan, Y. Liu, F. Liu, D. Yogatama, Y. Wang, L. Kong, and N. Collier, “Language models can see: plugging visual controls in text generation,” *arXiv preprint arXiv:2205.02655*, 2022.
- [24] C. Snell, S. Yang, J. Fu, Y. Su, and S. Levine, “Context-aware language modeling for goal-oriented dialogue systems,” *arXiv preprint arXiv:2204.10198*, 2022.
- [25] S. Verma, J. Fu, M. Yang, and S. Levine, “Chai: A chatbot ai for task-oriented dialogue with offline reinforcement learning,” *arXiv preprint arXiv:2204.08426*, 2022.
- [26] X. Chen, X. Wang, S. Changpinyo, A. Piergiovanni, P. Padlewski, D. Salz, S. Goodman, A. Grycner, B. Mustafa, L. Beyer, *et al.*, “Pali: A jointly-scaled multilingual language-image model,” *arXiv preprint arXiv:2209.06794*, 2022.
- [27] L. H. Li, M. Yatskar, D. Yin, C.-J. Hsieh, and K.-W. Chang, “Visualbert: A simple and performant baseline for vision and language,” *arXiv preprint arXiv:1908.03557*, 2019.
- [28] C. Sun, A. Myers, C. Vondrick, K. Murphy, and C. Schmid, “Videobert: A joint model for video and language representation learning,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019.
- [29] J.-B. Alayrac, J. Donahue, P. Luc, A. Miech, I. Barr, Y. Hasson, K. Lenc, A. Mensch, K. Millican, M. Reynolds, *et al.*, “Flamingo: a visual language model for few-shot learning,” *arXiv preprint arXiv:2204.14198*, 2022.
- [30] A. Suglia, Q. Gao, J. Thomason, G. Thattai, and G. Sukhatme, “Embodied bert: A transformer model for embodied, language-guided visual task completion,” *arXiv preprint arXiv:2108.04927*, 2021.
- [31] A. Pashevich, C. Schmid, and C. Sun, “Episodic transformer for vision-and-language navigation,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021.
- [32] P. Sharma, A. Torralba, and J. Andreas, “Skill induction and planning with latent language,” *arXiv preprint arXiv:2110.01517*, 2021.
- [33] C. Lynch and P. Sermanet, “Grounding language in play,” 2020.
- [34] S. Nair, E. Mitchell, K. Chen, B. Ichter, S. Savarese, and C. Finn, “Learning language-conditioned robot behavior from offline data and crowd-sourced annotation,” in *Conference on Robot Learning*, pp. 1303–1315, PMLR, 2021.
- [35] F. Hill, S. Mokra, N. Wong, and T. Harley, “Human instruction-following with deep reinforcement learning via transfer-learning from text,” *arXiv preprint arXiv:2005.09382*, 2020.
- [36] R. Zellers, X. Lu, J. Hessel, Y. Yu, J. S. Park, J. Cao, A. Farhadi, and Y. Choi, “Merlot: Multimodal neural script knowledge models,” *Advances in Neural Information Processing Systems*, 2021.
- [37] M. Reid, Y. Yamada, and S. S. Gu, “Can wikipedia help offline reinforcement learning,” *arXiv preprint arXiv:2201.12122*, 2022.
- [38] S. Li, X. Puig, Y. Du, C. Wang, E. Akyurek, A. Torralba, J. Andreas, and I. Mordatch, “Pre-trained language models for interactive decision-making,” *arXiv preprint arXiv:2202.01771*, 2022.
- [39] W. Huang, P. Abbeel, D. Pathak, and I. Mordatch, “Language models as zero-shot planners: Extracting actionable knowledge for embodied agents,” in *International Conference on Machine Learning*, PMLR, 2022.
- [40] M. Ahn, A. Brohan, N. Brown, Y. Chebotar, O. Cortes, B. David, C. Finn, K. Gopalakrishnan, K. Hausman, A. Herzog, D. Ho, J. Hsu, J. Ibarz, B. Ichter, A. Irpan, E. Jang, R. J. Ruano, K. Jeffrey, S. Jesmonth, N. Joshi, R. Julian, D. Kalashnikov, Y. Kuang, K.-H. Lee, S. Levine, Y. Lu, L. Luu, C. Parada, P. Pastor, J. Quiambao, K. Rao, J. Rettinghouse, D. Reyes, P. Sermanet, N. Sievers, C. Tan, A. Toshev, V. Vanhoucke, F. Xia, T. Xiao, P. Xu, S. Xu, and M. Yan, “Do as i can and not as i say: Grounding language in robotic affordances,” in *arXiv preprint arXiv:2204.01691*, 2022.
- [41] A. Zeng, A. Wong, S. Welker, K. Choromanski, F. Tombari, A. Purohit, M. Ryoo, V. Sindhwani, J. Lee, V. Vanhoucke, *et al.*, “Socratic models:

- Composing zero-shot multimodal reasoning with language,” *arXiv preprint arXiv:2204.00598*, 2022.
- [42] B. Chen, F. Xia, B. Ichter, K. Rao, K. Gopalakrishnan, M. S. Ryoo, A. Stone, and D. Kappler, “Open-vocabulary queryable scene representations for real world planning,” *arXiv preprint arXiv:2209.09874*, 2022.
- [43] D. Shah, B. Osinski, B. Ichter, and S. Levine, “Lm-nav: Robotic navigation with large pre-trained models of language, vision, and action,” *arXiv preprint arXiv:2207.04429*, 2022.
- [44] W. Huang, F. Xia, T. Xiao, H. Chan, J. Liang, P. Florence, A. Zeng, J. Tompson, I. Mordatch, Y. Chebotar, P. Sermanet, N. Brown, T. Jackson, L. Luu, S. Levine, K. Hausman, and B. Ichter, “Inner monologue: Embodied reasoning through planning with language models,” in *arXiv preprint arXiv:2207.05608*, 2022.
- [45] J. Liang, W. Huang, F. Xia, P. Xu, K. Hausman, B. Ichter, P. Florence, and A. Zeng, “Code as policies: Language model programs for embodied control,” *arXiv preprint arXiv:2209.07753*, 2022.
- [46] I. Singh, V. Blukis, A. Mousavian, A. Goyal, D. Xu, J. Tremblay, D. Fox, J. Thomason, and A. Garg, “Progprompt: Generating situated robot task plans using large language models,” *arXiv preprint arXiv:2209.11302*, 2022.
- [47] O. Mees, J. Borja-Diaz, and W. Burgard, “Grounding language with visual affordances over unstructured data,” *arXiv preprint arXiv:2210.01911*, 2022.
- [48] C. Huang, O. Mees, A. Zeng, and W. Burgard, “Visual language maps for robot navigation,” *arXiv preprint arXiv:2210.05714*, 2022.
- [49] S. S. Raman, V. Cohen, E. Rosen, I. Idrees, D. Paulius, and S. Tellex, “Planning with large language models via corrective re-prompting,” *arXiv preprint arXiv:2211.09935*, 2022.
- [50] C. H. Song, J. Wu, C. Washington, B. M. Sadler, W.-L. Chao, and Y. Su, “Llm-planner: Few-shot grounded planning for embodied agents with large language models,” *arXiv preprint arXiv:2212.04088*, 2022.
- [51] L. P. Kaelbling and T. Lozano-Pérez, “Hierarchical planning in the now,” in *Workshops at the Twenty-Fourth AAAI Conference on Artificial Intelligence*, 2010.
- [52] R. E. Fikes and N. J. Nilsson, “Strips: A new approach to the application of theorem proving to problem solving,” *Artificial intelligence*, 1971.
- [53] M. Toussaint, “Logic-geometric programming: An optimization-based approach to combined task and motion planning,” in *Twenty-Fourth International Joint Conference on Artificial Intelligence*, 2015.
- [54] D. Xu, S. Nair, Y. Zhu, J. Gao, A. Garg, L. Fei-Fei, and S. Savarese, “Neural task programming: Learning to generalize across hierarchical tasks,” in *2018 IEEE International Conference on Robotics and Automation (ICRA)*, 2018.
- [55] N. Savinov, A. Dosovitskiy, and V. Koltun, “Semi-parametric topological memory for navigation,” *arXiv preprint arXiv:1803.00653*, 2018.
- [56] T. Silver, R. Chitnis, N. Kumar, W. McClinton, T. Lozano-Perez, L. P. Kaelbling, and J. Tenenbaum, “Inventing relational state and action abstractions for effective and efficient bilevel planning,” *arXiv preprint arXiv:2203.09634*, 2022.
- [57] C. R. Garrett, C. Paxton, T. Lozano-Pérez, L. P. Kaelbling, and D. Fox, “Online replanning in belief space for partially observable task and motion problems,” in *2020 IEEE International Conference on Robotics and Automation (ICRA)*, 2020.
- [58] B. Eysenbach, R. R. Salakhutdinov, and S. Levine, “Search on the replay buffer: Bridging planning and reinforcement learning,” *Advances in Neural Information Processing Systems*, 2019.
- [59] D.-A. Huang, S. Nair, D. Xu, Y. Zhu, A. Garg, L. Fei-Fei, S. Savarese, and J. C. Niebles, “Neural task graphs: Generalizing to unseen tasks from a single video demonstration,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019.
- [60] D. Xu, A. Mandlekar, R. Martín-Martín, Y. Zhu, S. Savarese, and L. Fei-Fei, “Deep affordance foresight: Planning through what can be done

- in the future,” in *2021 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 6206–6213, IEEE, 2021.
- [61] B. Ichter, P. Sermanet, and C. Lynch, “Broadly-exploring, local-policy trees for long-horizon task planning,” *Conference on Robot Learning (CoRL)*, 2021.
- [62] C. Agia, T. Migimatsu, J. Wu, and J. Bohg, “Taps: Task-agnostic policy sequencing,” *arXiv preprint arXiv:2210.12250*, 2022.
- [63] D. Shah, P. Xu, Y. Lu, T. Xiao, A. Toshev, S. Levine, and B. Ichter, “Value function spaces: Skill-centric state abstractions for long-horizon reasoning,” *ICLR*, 2022.
- [64] A. V. Nair, V. Pong, M. Dalal, S. Bahl, S. Lin, and S. Levine, “Visual reinforcement learning with imagined goals,” in *Advances in Neural Information Processing Systems*, 2018.
- [65] B. Wu, S. Nair, L. Fei-Fei, and C. Finn, “Example-driven model-based reinforcement learning for solving long-horizon visuomotor tasks,” *arXiv preprint arXiv:2109.10312*, 2021.
- [66] F. Xia, C. Li, R. Martín-Martín, O. Litany, A. Toshev, and S. Savarese, “Relmogen: Integrating motion generation in reinforcement learning for mobile manipulation,” in *2021 IEEE International Conference on Robotics and Automation (ICRA)*, 2021.
- [67] D. Driess, O. Oguz, J.-S. Ha, and M. Toussaint, “Deep visual heuristics: Learning feasibility of mixed-integer programs for manipulation planning,” in *2020 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 9563–9569, IEEE, 2020.
- [68] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in neural information processing systems*, vol. 30, 2017.
- [69] M. Chevalier-Boisvert, L. Willems, and S. Pal, “Minimalistic gridworld environment for gymnasium,” 2018.
- [70] A. Zeng, P. Florence, J. Tompson, S. Welker, J. Chien, M. Attarian, T. Armstrong, I. Krasin, D. Duong, V. Sindhwani, and J. Lee, “Transporter networks: Rearranging the visual world for robotic manipulation,” *Conference on Robot Learning (CoRL)*, 2020.
- [71] M. Shridhar, L. Manuelli, and D. Fox, “Cliport: What and where pathways for robotic manipulation,” in *Conference on Robot Learning*, pp. 894–906, PMLR, 2022.
- [72] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.
- [73] A. G. Barto and S. Mahadevan, “Recent advances in hierarchical reinforcement learning,” *Discrete Event Dynamic Systems*, vol. 13, p. 41–77, jan 2003.
- [74] J. Wei, X. Wang, D. Schuurmans, M. Bosma, E. Chi, Q. Le, and D. Zhou, “Chain of thought prompting elicits reasoning in large language models,” *arXiv preprint arXiv:2201.11903*, 2022.
- [75] L. Van der Maaten and G. Hinton, “Visualizing data using t-sne,” *Journal of machine learning research*, vol. 9, no. 11, 2008.
- [76] A. Brohan, N. Brown, J. Carbajal, Y. Chebotar, J. Dabis, C. Finn, K. Gopalakrishnan, K. Hausman, A. Herzog, J. Hsu, *et al.*, “Rt-1: Robotics transformer for real-world control at scale,” *arXiv preprint arXiv:2212.06817*, 2022.
- [77] S. Reed, K. Zolna, E. Parisotto, S. G. Colmenarejo, A. Novikov, G. Barth-Maron, M. Gimenez, Y. Sulsky, J. Kay, J. T. Springenberg, *et al.*, “A generalist agent,” *arXiv preprint arXiv:2205.06175*, 2022.
- [78] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. L. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, *et al.*, “Training language models to follow instructions with human feedback,” *arXiv preprint arXiv:2203.02155*, 2022.
- [79] A. Chowdhery, S. Narang, J. Devlin, M. Bosma, G. Mishra, A. Roberts, P. Barham, H. W. Chung, C. Sutton, S. Gehrmann, *et al.*, “Palm: Scaling language modeling with pathways,” *arXiv preprint arXiv:2204.02311*, 2022.

- [80] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning transferable visual models from natural language supervision,” in *Proceedings of the 38th International Conference on Machine Learning* (M. Meila and T. Zhang, eds.), vol. 139 of *Proceedings of Machine Learning Research*, pp. 8748–8763, PMLR, 18–24 Jul 2021.
- [81] E. Jang, A. Irpan, M. Khansari, D. Kappler, F. Ebert, C. Lynch, S. Levine, and C. Finn, “Bc-z: Zero-shot task generalization with robotic imitation learning,” in *Conference on Robot Learning*, pp. 991–1002, PMLR, 2021.
- [82] M. Shridhar, L. Manuelli, and D. Fox, “Perceiver-actor: A multi-task transformer for robotic manipulation,” in *Proceedings of the 6th Conference on Robot Learning (CoRL)*, 2022.
- [83] M. Minderer, A. Gritsenko, A. Stone, M. Neumann, D. Weissenborn, A. Dosovitskiy, A. Mahendran, A. Arnab, M. Dehghani, Z. Shen, *et al.*, “Simple open-vocabulary object detection with vision transformers,” *arXiv preprint arXiv:2205.06230*, 2022.

A. Appendix

A.1. Grounded Decoding Implementation Details

We study three different implementations of Grounded Decoding for each of the experimental domains. While each instantiation applied Grounded Decoding to long-horizon planning and behavior synthesis, different components including language models and grounded models are used in each domain, as seen in Table 5. Grounded models used in these domains include Affordance Functions (AF), Safety Functions (S), Preference Functions (P), and Open-Vocabulary Object Detectors (D).

	Tabletop Rearrangement (Sim)	MiniGrid 2D Maze (Sim)	Kitchen Mobile Manipulation (Real)
LLM	InstructGPT [78]	InstructGPT	InstructGPT + PaLM[79]
Primitives	CLIPort [71]	PPO [72]	RT-1 [76]
Grounded Models	AF + S + P	AF	D

Table 5: Comparison between different versions of Grounded Decoding implemented in three different environments.

A.2. Implementation Details of Simulated Tabletop Rearrangement

A.2.1. Tasks

There are a total of 20 tasks (templates of language instructions), listed in Table 6, grouped into three task category: Letters, Blocks&Bowls, and Box Packing. Three categories share a total of 57 objects. For Letters category, the goals are to rearrange the alphabetical letter objects such that they satisfy certain orders specified by the language instructions. At the beginning of each episode, task-relevant objects and a set of 1 to 3 randomly-sampled distractor objects (except for the Letters category) are initialized at random positions on the tabletop with fixed orientations. A minimum 15cm distance is enforced between any two objects to avoid collision and penetration at initialization. To allow for automatic evaluations, a binary reward function is defined for each task using ground-truth state of the objects. Furthermore, we implement scripted policies for each task to collect demonstrations for training the CLIPort baseline. For certain tasks, we also randomize the attributes mentioned in the given instructions, which can be found below:

- $\langle \text{word} \rangle$: hi, world, left, right, top, down, love, you
- $\langle \text{corner/side} \rangle$: left side, top left corner, top side, top right corner, bottom right corner, bottom side, bottom left corner

A.2.2. Low-level Primitives

We use CLIPort [71] as the low-level primitive that can be invoked by the LLM planner, as it shows promising results of generalization across free-form language instructions. Additionally, since the policy predicts per-pixel affordance, it can be repurposed to serve as grounded models for planning for long-horizon tasks, which we leverage in this work. The single primitive policy is trained on 50,000 pre-collected demonstrations, across 10 training tasks, where each demonstration contains 1) language instruction of the format “pick up [x] and place it on [y]”, 2) top-down RGB-D observation of the current scene, 3) expert pick location expressed as pixel coordinates, and 4) expert place location expressed as pixel coordinates. The expert actions are obtained by accessing ground-truth object pose in the simulator.

A.2.3. Language Model

We use InstructGPT [78] (text-davinci-002), accessed through [OpenAI API](#).

A.2.4. Grounding Functions

Affordance Grounding Function (AF). The affordance should be a function of the robot, the environment, and the underlying policy in order to accurately determine what is possible in the scene. Given scene

image s and partially-decoded open-ended language instruction ℓ , the primitive policy CLIPort predicts unnormalized logits over the pixel space $s_{\text{pick}}, s_{\text{place}} \in \mathbb{R}^{480 \times 640}$, respectively for the pick location and the place location. Therefore, for any observation s and instruction ℓ , we calculate the affordance as $p_{\text{AF}}(s, \ell) = \max_{(x,y) \in 480 \times 640} (s_{\text{pick}}(x, y) + s_{\text{place}}(x, y))$. The affordance grounding function is used for all tasks.

Safety Grounding Function (S). World grounding for embodiment may not only be affordances, it can also be crucial robotics functions like safety. While any safety function that can score a language instruction ℓ may be used, we implement a simple indicator to prevent the robot from interacting with knives and red boxes, which we use as hazardous objects in this domain: $p_{\text{S}}(s, \ell) = \frac{1}{Z}(1 - \mathbb{I}[\text{red or knife in } \ell]) + \frac{\epsilon}{Z}\mathbb{I}[\text{red or knife in } \ell]$, where Z is the normalizing term and ϵ is a small value for ensuring the joint probability does not collapse to 0. Safety grounding function is used for 3 tasks in Box Packing task family.

Preference Grounding Function (P). Robots operating alongside humans should also be aware of human preferences, which often differ based on the specific user. We choose two random objects (o_1, o_2) as the preferred objects to be used for the household object sorting tasks. Note that unlike safety functions, preferences often come in the form of “soft requirement”. Therefore, the preference grounding function is implemented as $p_{\text{P}}(s, \ell) = \frac{x}{Z}\mathbb{I}[o_1 \text{ or } o_2 \text{ in } \ell] + \frac{y}{Z}(1 - \mathbb{I}[o_1 \text{ or } o_2 \text{ in } \ell])$, where we choose $x = 0.5$ and $y = 0.1$ for our experiments. More generic preference functions may also be learned or statistically calculated based on history of user data. Preference grounding function is used for 2 tasks in Box Packing task family.

A.2.5. CLIPort Baseline

As CLIPort [71] already takes as input a natural language instruction and is capable of directly outputting robot actions, it bears the question whether we need a high-level planner for completing long-horizon tasks. To this end, we additionally train two variants of multi-task CLIPort policy on 10 of the total 20 tasks as baselines (see Table 6 for the train/test split). One variant, which we referred as “CLIPort (Short)”, is trained only on single-step pick-and-place instructions of the format “pick up [x] and place it on [y]” on the 10 training tasks. The decomposed pick-and-place instructions are obtained from scripted planners. At evaluation time, the policy is fed in only the high-level instructions without any planners. The other variant, which we referred as “CLIPort (Long)”, is trained on the high-level instructions from the 10 training tasks (without decomposition from scripted planners). Similarly, at evaluation time, it is fed in only the high-level instructions and evaluated on both seen and unseen instructions. Both variants are trained on 50,000 demonstrations, similar to the Grounded Decoding primitive. The goal of these baselines is to evaluate whether solitary language-conditioned policies can perform well on long-horizon tasks and generalize to new task instructions. Note that the CLIPort baselines are different from the primitive used in Grounded Decoding, although they share the same architecture.

A.2.6. Full Experimental Results in Simulated Tabletop Domain

Below we show the full list of tasks and the full experimental results in the simulated tabletop domain. Each entry is the average success rate across 20 rollouts. The tasks with blue-colored background are seen tasks and the tasks with orange-colored background are the unseen tasks. Seen tasks may be used for training for supervised baselines (CLIPort) or may be used in the prompt for methods using language model planner. Note that for the “Box Packing” task category, although all tasks were seen in training or the prompts, we enforce additional safety and preference constraints for evaluation only at test time.

Tasks	p_G	CLIPort		+LLM	+Grounded Decoding	
		Short	Long	Ungrounded	Greedy	Beam
Letters						
Put the letters in alphabetical order from left to right	AF	5%	20%	10%	20%	40%
Spell as much of <i>word</i> as you can	AF	10%	60%	30%	60%	65%
Separate the vowels from the remaining letters to the bottom side	AF	5%	40%	20%	50%	65%
Put the letters in reverse alphabetical order from left to right	AF	15%	10%	15%	25%	25%
Correctly spell out a sport using the present letters	AF	10%	10%	5%	30%	30%
Sort the geometrically symmetrical letters to the bottom side	AF	5%	10%	15%	35%	50%
Separate the consonants from the remaining letters to the bottom side	AF	0%	0%	25%	25%	25%
Sort the letters less than "D" according to ASCII to the bottom side	AF	0%	20%	35%	70%	75%
Blocks & Bowls						
Stack all the blocks	AF	5%	90%	30%	75%	90%
Put all the blocks on the <i>corner/side</i>	AF	0%	65%	50%	45%	70%
Put all the blocks in the bowls with matching colors	AF	0%	30%	25%	60%	70%
Put the blocks in the bowls with mismatched colors	AF	25%	30%	45%	30%	55%
Put all the blocks in different corners	AF	0%	5%	40%	50%	60%
Stack only the blocks of cool colors	AF	5%	5%	20%	70%	70%
Stack only the blocks of warm colors	AF	0%	10%	15%	45%	35%
Sort the primary color blocks to the left side	AF	0%	0%	20%	25%	30%
Box Packing						
Pack the objects into the brown box	AF + S	20%	40%	5%	100%	90%
Pack the objects into the boxes	AF + S	10%	20%	5%	75%	70%
I'd like some snacks on the right side	AF + P	15%	20%	15%	40%	55%
Pack me a picnic box	AF + S + P	15%	30%	20%	100%	95%

Table 6: Full Experimental Results in Simulated Tabletop Rearrangement Tasks. The tasks with blue-colored background are seen tasks and the tasks with orange-colored background are the unseen tasks. *Box Packing tasks are all seen during training, but safety and preference requirements are only enforced during evaluation.

A.3. Implementation Details of MiniGrid 2D Maze

A.3.1. Environment Setup

We use the open-source `gym-minigrid` suite of environments to evaluate our method with one simple change — instead of the default observation space which is a 7×7 egocentric window, our agent has access to *entire grid* — that allows us to simplify the tasks by removing partial observability [63].

A.3.2. Tasks

The tasks are grouped in three categories (please see Table 7 for example instructions):

1. **Easy:** Simple tasks where the horizon is short (10-30 steps) and fully described by the textual instruction, e.g. `OpenDoors` and `PutNext`. The short horizon makes them relatively easy for a wide range of HRL algorithms. The instructions for these tasks generally spell out each individual skill, making them particularly easy for high-level planners based on language modeling.
2. **Medium:** Combination of short and long horizon tasks (up to 80 steps) with step-by-step textual instructions, e.g. `LockedRoom`. While being significantly longer, these tasks also tend to have instructions that spell out the low-level tasks (see Table 7).
3. **Hard:** Complex, long horizon instructions (over 100 steps) with short, ambiguous instructions that necessitate multi-step reasoning and efficient exploration, e.g. `MultiRoom` and `BlockedUnlock`. In addition to being long-horizon, the instructions in this case tend to be ambiguous and under-specified, e.g. "traverse through the rooms to get to the goal", which does not provide enough context for any *blind* planning agent.

Difficulty	Task Name	Example Instruction
Easy	<code>OpenDoors</code>	open door blue, then open door red
	<code>PutNext</code>	move the red ball next to the green box
Medium	<code>LockedRoom</code>	get the red key from the purple room, open the red door and go to the goal
Hard	<code>MultiRoom</code>	traverse the rooms to get to the goal
	<code>BlockedUnlock</code>	pick up the blue box

Table 7: Example Instructions in Minigrid

A.3.3. Language Model

We use InstructGPT [78] (`text-davinci-002`), accessed through [OpenAI API](#). The prompts used can be found in Section A.5.

We found the prompts to be generally sufficient for solving the “seen” tasks, as well as “unseen” tasks, i.e. tasks that do not have an example in the context. Empirically, we did not find any improvements by including more than 3 example tasks in the prompt — we hypothesize that this is likely due to the shared low-level primitives across tasks. For all Minigrid experiments presented in this paper, we used the prompt shown in Section A.5.

A.3.4. Low-level Primitives

To train low-level primitives, we train an RL agent to solve a wide range of short-horizon subtasks (under 10 steps) that are shared across the various Minigrid tasks — `go to <obj>`, `pick up <obj>`, `drop <obj>`, `open <obj>`. Rather than training individual skills for each of them [63], we train a single multi-task policy that is conditioned on the CLIP embeddings [80] of the task strings. This scheme allows some robustness to synonyms and ambiguous task specifications, and has been widely used in learning language-grounded policies [81, 82].

We train these primitives using PPO [72], as recommended by the environment developers [69]. Each of these skills are trained with a sparse outcome reward (+1 if a trajectory is successful, 0 otherwise). In addition to these low-level skills, we perform a form of hindsight relabeling where “substrings” of the task strings are masked to allow generalization to partial strings, e.g. “go to red” may be interpreted as “go to red key” or “go to red door”, and our masking strategy allows the multi-task policy to execute tasks specified by partially complete strings, if necessary.

A.3.5. Grounded Affordance Function

We use the task string-conditioned value function estimates from our learned policy to obtain a visually grounded affordance function.

A.3.6. Additional Qualitative Results

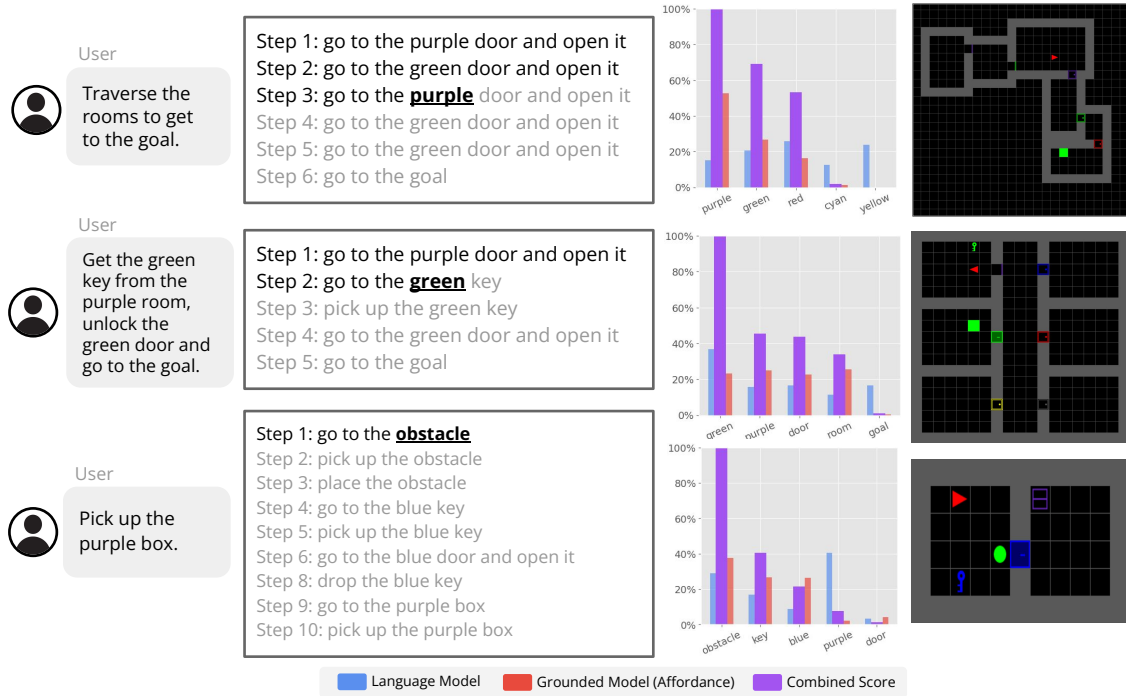


Figure 8: Minigrid Domain

A.4. Implementation Details of Real-World Mobile Manipulation

A.4.1. Tasks

Instruction
put an energy bar and water bottle on the table
bring me a lime soda and a bag of chips
Can you throw away the apple and bring me a coke
bring me a 7up can and a tea
move an multigrain chips to the table and an apple to the far counter
move the lime soda, the sponge, and the water bottle to the table
bring me an apple, a coke, and water bottle

Table 8: List of unambiguous SayCan instructions

Instruction
I want to wipe some spill.
Bring me a fruit
Bring me a snack
Bring me a bag of chips
Bring me a bag of snack
Bring me a bag of chips and something to wipe a spill
Bring me a bag of chips and something to drink
Bring me a bag of chips and a soda
Human: I want a soda that is not coke, and a fruit.
I want a fruit and a soda

Table 9: List of ambiguous SayCan instructions

A.4.2. Language Model

For planning, we use PaLM [79], a 540B parameter language model trained on a large datasets that include high-quality web documents, books, Wikipedia, conversations, and GitHub code. Before planning, we use InstructGPT [78] (text-davinci-002), accessed through OpenAI API. to generate the (grounded) chain of thought.

We used square bracket to indicate grounded decoding, as illustrated in Fig. 5. The prompts are shown in Listing 3.

A.4.3. Low-level Primitives

We use a combination of learned and scripted control policies for navigation and manipulation, following the implementation described in SayCan [40] and RT-1 [76]. The manipulation policies for the picking action are learned using Behavior Cloning (BC) on 68000 demonstrations and 12000 autonomous successes that were collected over the course of 11 months using a fleet of 10 robots. The demonstrations are collected by teleoperators using VR headset controllers to track the motion of their hand, which is then mapped onto the robot’s end-effector pose. The navigation policies are scripted, based on a ground-truth map as well as a learned perception module for collision avoidance and planning. The placing actions follow pre-computed motions only when preceded by a navigation policy. The Value Functions used by SayCan for affordance

grounding are provided by the Q -networks of trained RL agents; we follow the RL training setup described in [40].

A.4.4. Open-Vocabulary Detector Grounding Function

We use owl-vit [83] as our grounding model. It takes in an image and a natural language query, and returns a list of bounding boxes with scores. We take the maximum score as the grounding function.

More examples of object detection as a grounding function can be found in Fig. 9.

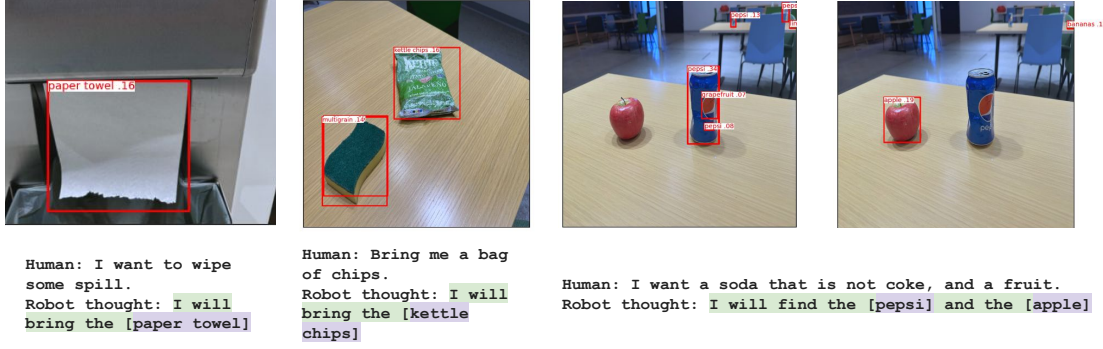


Figure 9: Additional examples of using open-vocabulary object detection as a grounding function in Real-World Kitchen Mobile Manipulation Domain.

A.5. Prompts

Listing 1: Grounded Decoding Prompt in Simulated Tabletop Rearrangement Domain

```

Task: Pack all letter objects on the brown box
Step 1: pick up the e and place it on the brown box
Step 2: pick up the g and place it on the brown box
Step 3: done

Task: Put the letters on the tables in alphabetical order
Step 1: pick up the c and place it on the bottom left side
Step 2: pick up the d and place it on the right of c
Step 3: pick up the i and place it on the right of d
Step 4: pick up the l and place it on the right of i
Step 5: pick up the w and place it on the right of l
Step 6: done

Task: Spell as much of "blue" as you can
Step 1: pick up the l and place it on the bottom left side
Step 2: pick up the the u and place it on the right of l
Step 3: pick up the the e and place it on the right of u
Step 4: done

Task: Separate the vowels from the remaining letters
Step 1: pick up the i and place it on the bottom side
Step 2: pick up the o and place it on the bottom side
Step 3: done

Task: Stack all the blocks
Step 1: pick up the brown block and place it on the pink block
Step 2: pick up the cyan block and place it on the brown block
Step 3: pick up the orange block and place it on the cyan block
Step 4: pick up the gray block and place it on the orange block
Step 5: done

Task: Put all the blocks on the bottom left corner
Step 1: pick up the white block and place it on the bottom left corner
Step 2: pick up the yellow block and place it on the bottom left corner
Step 3: pick up the green block and place it on the bottom left corner
Step 4: pick up the blue block and place it on the bottom left corner
Step 5: pick up the purple block and place it on the bottom left corner
Step 6: done

Task: Put all the blocks in the bowls with matching colors
Step 1: pick up the cyan block and place it on the cyan bowl
Step 2: pick up the purple block and place it on the purple bowl
Step 3: pick up the brown block and place it on the brown bowl
Step 4: pick up the pink block and place it on the pink bowl
Step 5: done

Task: Pack the items into any box
Step 1: pick up the donut stick and place it on the red box
Step 2: pick up the pepsi and place it on the brown box
Step 3: pick up the peach and place it on the brown box
Step 4: pick up the strawberry and place it on the red box
Step 5: done

Task: Pack the items on the table into the brown box
Step 1: pick up the knife and place it on the brown box
Step 2: pick up the plum and place it on the brown box
Step 3: pick up the pepsi and place it on the brown box
Step 4: pick up the cupcake and place it on the brown box
Step 5: done

Task: Pack the items on the table into the brown box
Step 1: pick up the i and place it on the brown box
Step 2: pick up the green block and place it on the brown box
Step 3: pick up the l and place it on the brown box
Step 4: done

Task: Can you put some snacks on the right side for me?
Step 1: pick up the plum and place it on the right side

```

Step 2: done

Task: Can you pack my picnic box for me?

Step 1: pick up the orange and place it on the picnic box

Step 2: pick up the diet pepsi and place it on the picnic box

Step 3: pick up the knife and place it on the picnic box

Step 4: done

Listing 2: Grounded Decoding Prompt in MiniGrid 2D Maze Domain

You are a 2D maze-solving agent with access to a variety of low-level skills such as picking up or dropping objects, navigating to doors/keys/boxes, and opening/closing doors. Here are some example tasks:

Task: get the green key from the purple room, unlock the green door and go to the goal

Step 1: go to the purple door and open it

Step 2: go to the green key

Step 3: pick up the key

Step 4: go to the green door and open it

Step 5: go to the goal.

Task: pick up the purple box

Step 1: go to the green obstacle

Step 2: pick up the obstacle

Step 3: place the obstacle

Step 4: go to the blue key

Step 5: pick up the blue key

Step 6: go to the blue door and open it

Step 8: drop the blue key

Step 7: go to the purple box

Step 8: pick up the purple box.

Task: traverse the rooms to get to the goal

Step 1: go to the purple door and open it

Step 2: go to the green door and open it

Step 3: go to the purple door and open it

Step 4: go to the green door and open it

Step 5: go to the green door and open it

Step 6: go to the goal.

Now your turn.

Listing 3: Grounded Decoding Prompt in Real-World Kitchen Mobile Manipulation Domain

The following objects are in the scene: 7up, apple, banana, mango, tea, multigrain chips, kettle chips, jalapeno chips, rice chips, coke, grapefruit soda, pepsi, redbull, energy bar, lime soda, sponge, paper towel, and water bottle.
The following locations are in the scene: close counter, far counter, table, trash, bowl.
The robot will always put object name in brackets [].

Robot: I am a robot that can bring objects to you.

Human: I am hungry.

Robot thought: I will find the [multigrain chips].

Robot plan: 1. Find the multigrain chips
2. Pick up the multigrain chips
3. Bring it to you
4. Put it down
5. Done

Robot: I am a robot that can bring objects to you.

Human: Throw away the fruit.

Robot thought: I will find the [mango] and move it to the trash.

Robot plan: 1. Find the mango
2. Pick up the mango
3. Go to the trash
4. Put it down
5. Done

Robot: I am a robot that can bring objects to you.

Human: (inject instruction).

Robot thought: